



École Polytechnique Fédérale de Lausanne

Unsupervised deep-learning-based video registration

by Luca Barras

Bachelor Project Report

Approved by the Examining Committee:

Prof. S. Süsstrunk
Thesis Advisor

Hakki Can Karaimer
Thesis Supervisor

EPFL IC IINFCOM IVRL
(Bâtiment BC)
Station 14
CH-1015 Lausanne

June 5, 2020

Abstract

This project aims to align an Infrared (IR) and a RGB video using an unsupervised deep-learning method. It exist already methods to align a pair of IR and RGB image but not again on videos. Because a video is composed of thousands of images, the previous methods was not adapted for as many images. The novelty is that we use an unsupervised deep-learning method to resolve this problem. Receiving a pair of RGB and IR video, we convert the video in frames and we send them to a convolutional neural network with a spatial transform layer. The spatial transform layer allows to transform the frames by applying linear or not linear transformation on the location of the pixels. The output of this network is the IR frame after the transformation that is more aligned with the RGB frame. We proceeded in three steps. First we used the gray-scale version of the frames as input of the network, then the contours of the frames and finally the human segmentation of the frames. After to have all the output frames we can reconstruct the video by merging the RGB frames and the new IR frames by applying an alpha value on them. The source code of this project is available online at https://github.com/lbarras27/deeplearning_registration_IR_RGB_video.

Contents

Abstract	2
1 Introduction	5
1.1 Setting	5
1.2 Goals and Motivation	5
2 Background	7
2.1 Infrared light	7
2.2 U-Net Network	7
2.3 Spatial Transformer Network	8
2.4 Voxelmorph	9
2.4.1 Architecture	9
2.4.2 Loss	10
3 Method	11
3.1 Voxelmorph network	11
3.1.1 Network architecture	11
3.1.2 Loss function	12
3.2 Affine network	13
3.2.1 Network architecture	13
3.2.2 Loss function	14
3.3 Transformation on the original frames	14
4 Experiments	15
4.1 Dataset	15
4.2 Implementation	17
4.3 Results	18
4.3.1 Gray-scale version	18
4.3.2 Contours version	19
4.3.3 Human segmentation version	20
5 Discussion	21
5.1 Future work	22

6 Conclusion	24
Bibliography	25

Chapter 1

Introduction

1.1 Setting

Nowadays the infrared (IR) is used in several domains, in scientific, military, industrial, law enforcement and medical applications. The most known use is for the night vision. Indeed the infrared cameras capture the heat in contrary to the usual cameras that capture the color that we see. In many applications we want to be able to align an ordinary image (RGB image) with an IR image, but it take works. For example in the medical domain, it is interesting to superpose the two images or the two videos to better see where is the anomalies. If to align an image can be laborious, imagine you to align a video that contain thousands of images in general.

On other side, the deep-learning is becoming more and more used in the scientific domain, mainly in the computer vision domain. For example, it is used for facial recognition or also in the autonomous car. Born in the 80s, the deep learning become really used in the 2010s thanks to the GPU use for this task. The deep learning is probably the most powerful domain in machine learning. The machine learning principle is to train a machine (program) to perform a certain task by giving it a lot of data in input.

1.2 Goals and Motivation

Why do not use the powerful of deep learning to align an IR video with a RGB video ? The alignment of images or videos is called the registration in the jargon of the computer vision engineers. It is that this project aims. We use an unsupervised deep-learning method to achieve this task. Using a supervised method is very laborious because we must label every frames of the video, so unsupervised method seem to be better for this task. There have already been studies on the registration of two IR, RGB images but never again with an unsupervised deep

learning method. Recently, there have been a method for the registration of two images with an unsupervised deep-learning network but not with IR images. This project will aims to use this method on IR and RGB videos as dataset. The dataset is a series of IR videos with their corresponding RGB videos. The IR and RGB videos are nearly the same but with a small shift because it is really complicated or even impossible to have exactly the same framing with two different cameras.

Chapter 2

Background

2.1 Infrared light

This is an electromagnetic radiation that have wavelengths longer than those of visible light. So it is invisible to the human eye. For a human the light is visible at a wavelength between 400nm and 700nm but the infrared light has a wavelength between 700nm and 1mm. Almost all heat sources emit infrared light. It is why the infrared camera can see in the night because it capture the heat. We can see on the figure 2.1 that the tea-pot contain a warm liquid for example.

2.2 U-Net Network

The U-Net Network is a special Convolution Neural Network (CNN) that it takes an input and the output has an higher definition than the classical CNN. For example, if we have an image of 300x200 as input, the output can be of same size, 300x200. This network allows to have a higher definition output.

On the figure 2.2 we can see a typical U-Net network architecture. This network is separate in




Figure 2.1: Infrared image and ordinary image (RGB)




Figure 2.2: U-Net network architecture, figure taken from [6]

two parts. The first one is the encoder, in this part the size of the features maps is reduce. The second part is called decoder because the size of the features maps increase. We remark we have connection between each same level where we copy the features maps from the encoder to the decoder. In the encoder part, in general there are a series of convolutions with stride greater than 1 or convolution with pooling (for example max-pooling). In general, we increase the number of channels when we go down and in the decoder part, on the contrary we do convolutions with up-sampling operations and we decrease the number of channels.

2.3 Spatial Transformer Network

A spatial transformer network (STN) is a generalization of differentiable attention to any spatial transformation. It allow a neural network to learn how to perform spatial transformations on the input image in order to enhance the geometric invariance of the model. It can crop, scale, translate and rotate a image for example. It is useful because the convolutional neural networks (CNN) are not invariant to affine transformations in general. Affine transformations are mainly the rotation, scale and translation.

STN is separate in three parts. We can see on the figure 2.3. The first part is the localization network. It is a CNN which regresses the transformation parameters. The second part is the grid generator, it generate a grid that map the position of each pixel of the input image to the corresponding position in the output image. We can see that more precisely on the figure 2.4. And the third part, called the sampler, use the coordinate of the grid to map the pixel positions of the input image to the output image (image after transformation).




Figure 2.3: Spatial Transformer architecture, figure taken from [7]




Figure 2.4: In (a), we can see an identity transformation. In (b), we have an affine transformation. U is the initial image and V is the image after the transformation, figure taken from [7]

2.4 Voxelmorph

Voxelmorph is a deep-learning based registration library write in python with a Keras and Pytorch version. To achieve to align two 2D or 3D images, it use an unsupervised deep-learning method.

2.4.1 Architecture

The network take two inputs, the fixed image and the moving image. Firstly, it use a U-Net network to get the registration field and then it use a spatial transform on the moving image with the registration field as transformation parameter. The output is the moved image that is the moving image after to have apply the transformation on it. The figure 2.5 show more precisely how is this network.

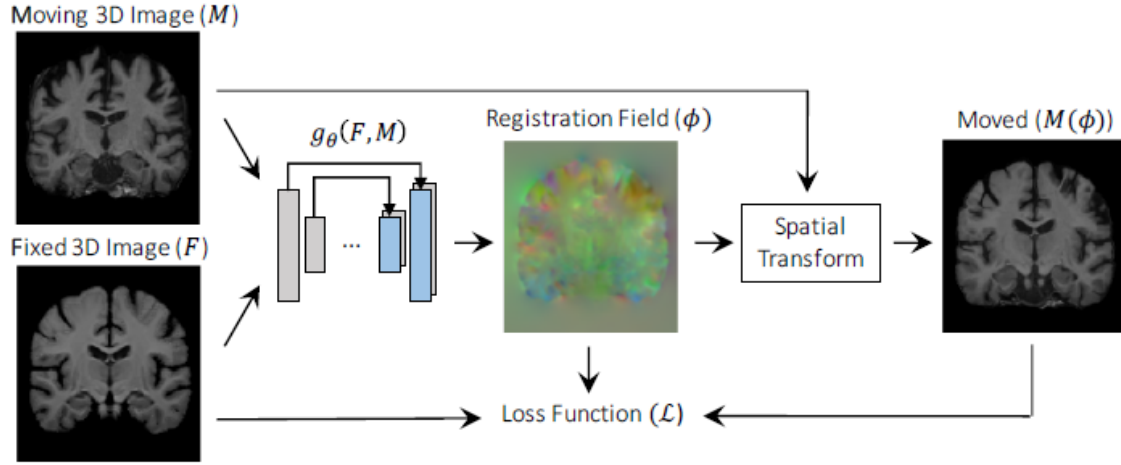


Figure 2.5: Voxelmorph architecture, figure taken from [4]

2.4.2 Loss

The loss function of this network is:

$$\mathcal{L}(F, M, \phi) = \mathcal{L}_{sim}(F, M(\phi)) + \lambda \mathcal{L}_{smooth}(\phi)$$

$\mathcal{L}_{sim}(\cdot, \cdot)$ measure the similarity between the fixed input and the moved image. For example $\mathcal{L}_{sim}(\cdot, \cdot)$ can be the mean square error between these two images. $\mathcal{L}_{smooth}(\phi)$ allow the regularization on ϕ . This part impose the smoothness of the transformation. λ is the regularization parameter.

Chapter 3

Method

To achieve the registration between the IR video and the RGB video, we will use two similar networks. First we will use the exactly same network like in Voxelmorph. So we have two inputs in the convolution neural network, the fixed image and the moving image and two output, the flow field and the moved image. In our case, the fixed input is a frame of the RGB video and the moving input is the corresponding frame in the IR video. On the figure 2.5 like explained before in the Voxelmorph architecture section.

For the second network, we use an architecture very similar to the previous one but there is a little difference. We start with a CNN to regress to a 2×3 affine matrix. And apply this matrix on the image. So the shape in the image cannot be deform because we apply an affine matrix on the image. In the previous network, we apply a flow field in the spatial transform component. So the pixels can move anywhere, it is for that we add the $\mathcal{L}_{smooth}(\phi)$ part in the loss function. The second network do not need this part in the loss function because the linearity of the transformation is guaranteed by the affine matrix.

We are going to proceed in three steps, at each step we will try with the two networks. The first step, we will just use the gray-scale version of the RGB and IR frames as input of the networks. In the second step, we will use the contours of the RGB and IR frames and in the last step, we will use the human segmentation of the frames as input of the networks.

3.1 Voxelmorph network

3.1.1 Network architecture

It is the same network used in Voxelmorph. It take as input the concatenation of the RGB frame and the IR frame. So the size is of $2 \times 320 \times 240$ in our case. It begin to enter in the U-Net network.

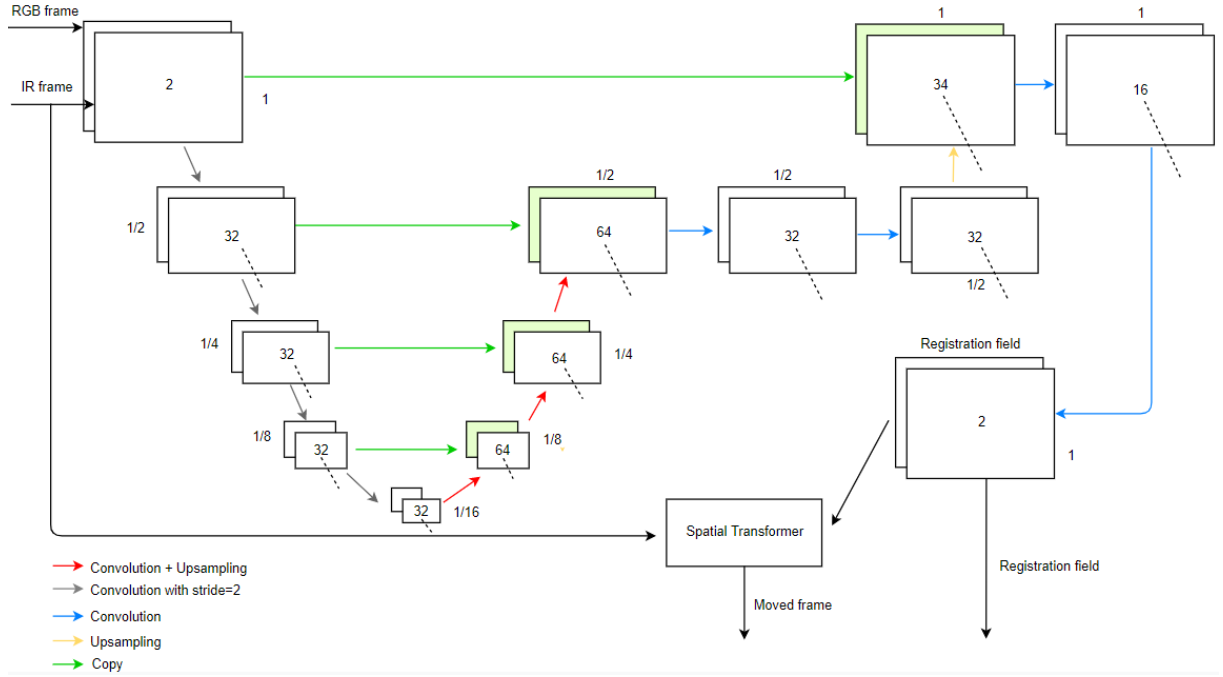


Figure 3.1: The number in center of each box is the number of channels. The number next to or above the box is the spatial resolution with respect to the input. The green boxes are the copied features maps. After each convolution, there is a Leaky ReLU activation.

In the first part, it does a series of 2D convolutions with stride equal to 2 followed by a leaky ReLU activation. So the size of the feature maps decrease layer by layer by a factor 2. In the second part, it does 2D convolutions with stride equal to 1 followed by a Leaky ReLU and an up-sampling of factor 2 and a concatenation with the feature map from the encoder which is on the same level. The kernels size of each convolution is 3 x 3. The output of this U-Net network give us the flow field to pass to the spatial transformer. The second input of the spatial transformer is the IR image (moving image). The function in the spatial transformer is simply the addition of the flow field with an identity grid. So the result grid give us the position to each pixel to map from the original IR frame to the new frame. The output of the spatial transformer is the IR image after the transformation. The two outputs of this network is firstly, the moved image of size 320 x 240 and secondly the flow fields of size 2 x 320 x 240. The network is showed in figure 3.1.

3.1.2 Loss function

The loss function is the same used in voxelmorph. Remember you, it was:

$$\mathcal{L}(F, M, \phi) = \mathcal{L}_{sim}(F, M(\phi)) + \lambda \mathcal{L}_{smooth}(\phi)$$

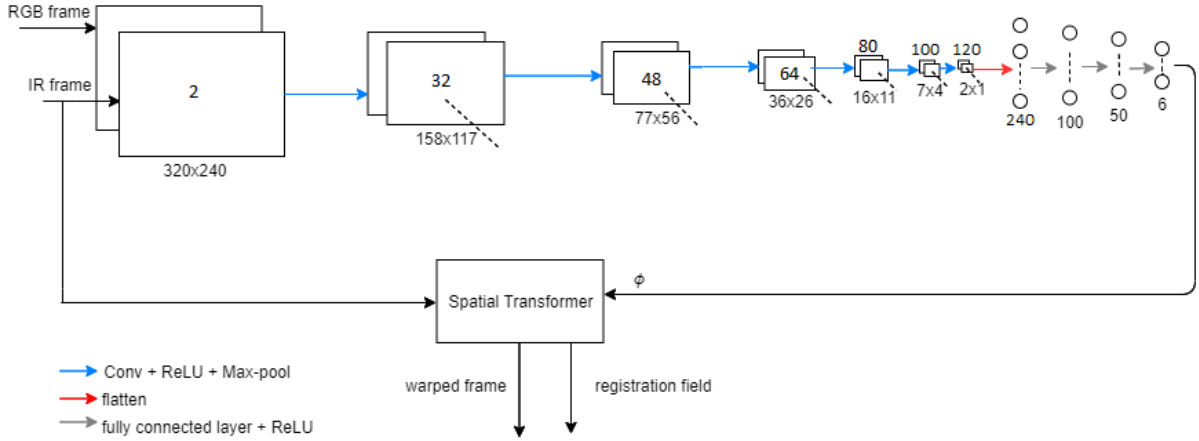


Figure 3.2: The number in center of or above each box is the number of channels and the numbers underneath the boxes are the size of the features maps.

Where F is the RGB image, M is the IR image and $M(\phi)$ is the network output. $\mathcal{L}_{sim}(F, M(\phi))$ measure the similitude between the IR image and the warped image. In this experiment we use the mean square error as function. $\mathcal{L}_{smooth}(\phi)$ force the smoothness of the transformation. In our case, $\mathcal{L}_{smooth}(\phi) = \|\nabla \phi\|^2$. Finally, λ is the regularization parameter, bigger it is, better is the smoothness of the transformation.

3.2 Affine network

3.2.1 Network architecture

The network is simply a spatial transformer network. It take as input an 2-channel image that is the concatenation of the IR frame and the RGB frame. The input size is of 2 x 320 x 240 in our case. The network begin with a succession of 2D convolutions followed by ReLU activation with kernel size of 7 x 7, 5 x 5, 5 x 5, 5 x 5, 3 x 3, 3 x 3 respectively. All the convolutions are did with a stride of size 1. After we flatten the size to use the fully connection layers. In our experiment, after the flattening, we have a size of 240 x 1. Then the next layers have a size of 100, 50 and 6 respectively. 6 can be seen as 2 x 3 that is the size of the affine matrix. In the spatial transformer we apply the affine matrix on the IR image and get the first output (warped image) of the network. The second output is the flow field generate by the affine matrix. The first output has a size of 320 x 240 and the second a size of 2 x 320 x 240. 2 channels because we have a channel for each direction (x, y). We can see more in details the network architecture on the figure 3.2.

3.2.2 Loss function

The loss function is simply the similitude between the RGB image (fixed input) and the warped image. It can be expressed as:

$$\mathcal{L}(F, M, \phi) = \mathcal{L}_{sim}(F, M(\phi))$$

Where F is the RGB image, M is the IR image and $M(\phi)$ is the IR image after to be went out of the spatial transformer. \mathcal{L}_{sim} is the mean square error in our experiment.

3.3 Transformation on the original frames

Because we do not use the original frames as input of the network but the gray-scale in the first case, the contours in the second case and the human segmentation in the last case, we must use the flow field output of the network to apply the transformation on the original frames (IR frames). When we have this flow field, we just need to put each channel (r, g, b) of the original frame in the spatial transformer and we obtain the IR moved frame that must be better align to the RGB frame.

Chapter 4

Experiments

4.1 Dataset

The dataset is a series of IR and RGB video pairs. As our dataset is a series of videos, we must do some preprocessing on it before to can send it in our network. Firstly, we focus on one video pair. The RGB and IR video last 21 minutes each. Then we convert the videos in frames and we obtain 16441 frames for each video. When we have the frames of each video, we must resize the infrared frames because they do not have the same size as the RGB frames. Indeed the IR frames has a size of 320 x 256 pixels and the size of the RGB frames are of 320 x 240 pixels. So we resize the IR frames in 320 x 240 pixels. After that, we remark that the RGB frames contain a lot of very dark images. We remove all these dark frames and the respectively IR frames because we always must have the pair of image to can use in our network. After this operation, we only have 12633 frames. Before to put these data in the network, we must again normalize the pixels value between 0 and 1. Now the data is ready to be send in the network.

In the first part of the experience, we just use the gray-scale version of the IR and RGB images.

In the second part, we want to work with the contours of the images. So we apply some filters to the images pairs to get their contours. Before each filters to get the contour of the images, we apply a Gaussian filter to smooth a little bit the images and get better results. The first filter we



Figure 4.1: Example of a frame with too much blue.




Figure 4.2: Example of pairs of IR and RGB frames after to have applying filters. On the top left with Sobel filter, top right with Canny filter, bottom left with adaptive threshold filter and bottom right with Laplacian filter.

apply is a Sobel filter in each direction (x, y) and compute its magnitude: $\sqrt{(sobel_x)^2 + (sobel_y)^2}$. The second filter is the Canny filter, the third one adaptive threshold filter and the last one laplacian filter. We can see the results of each filter on the figure 4.2.

In the last part, we use the human segmentation of the images. If the pixel belongs to a human, it has 0 as pixel value otherwise 255. So in this part we use only frames where there are humans. For the RGB images, we use a pretrain deep network call Deeplab to segment the images. We also try to use an another deep network call mask-RCNN but it is slower and give similar results. For the IR images, we threshold them with their blue channel. We must adapt the threshold following the frames, because some frames are lighter than others. In this part, we use a different dataset than previously because we must keep only frames where there are humans and select only the best results of the segmentation. We can see in figure 4.1, some IR frames are too difficult to threshold, because the frames are too blue. So to build this dataset, we proceeded as following. First we selected on each video only good IR frames we can easily threshold (humans having enough contrast with their environment). We keep too the associated RGB frames. Then we threshold the IR frames by sequences of frames, varying the threshold following the blue intensity of the sequence. After that we segment the associated RGB frame with Deeplab. Because sometimes Deeplab not detect humans, we must remove all the white RGB segmented frames with their respective IR segmented, RGB and IR frames. And when we did that, we must again check and remove all the bad frames of the RGB segmented frames. So finally we have 10431 4-tuple (RGB, IR, mask RGB, mask IR frames). We split it in 8000 frames for the training set, 1000 frames for the validation set and 1400 frames for the test set.



Figure 4.3: Example of a pair of IR and RGB frames after segmentation.




Figure 4.4: Example of a result of the print function on the masks

4.2 Implementation

To manipulate the dataset we use OpenCV and Numpy libraries. The networks are coded using Pytorch. For the first network (Voxelmorph network), we use the Pytorch version of the network implemented in the Voxelmorph paper but with little modifications. Indeed, in their Pytorch version, sometimes the code is not adapted to work with 2D images but only for 3D images. So we rewrite their network to make it work with 2D images. For the two networks, we use ADAM optimizer with a learning rate of 0.0001. To speed up the learning, we use a batch of 20 pairs of images for the two networks. To segment the RGB frames, we use Deeplabv3 pretrain on Pascal Voc dataset. To better see if an IR frame is aligned with its corresponding RGB frame, we implemented a print function that show the IR image, RGB image, warped image, the difference between the RGB and IR image and the difference between the RGB and the warped image. Figure 4.4 show us a typical print of this function. This function is useful when we work with the mask version of the images. The second tool we implemented to see if the frames are aligned is to reconstruct the video after to have merged the warped frame (IR frame after the transformation) and RGB frame. To merge the two images, we apply an alpha blending on them. Alpha blending is just an interpolation between the two images. The formula is given by: $\alpha M(\phi) + (1 - \alpha)F$, where F is the RGB frame and $M(\phi)$ is the warped frame. Figure 4.5 show the result of the merging of the two images with $\alpha = 0.35$. We also implemented all kind of functions to facilitate the manipulation of the dataset, for example to load, save, threshold and so on. To see more in details, check the source code of this project, it is available online at https://github.com/lbarras27/deeplearning_registration_IR_RGB_video.



Figure 4.5: We can see the result of the merge of the RGB and IR frame with $\alpha = 0.35$.

4.3 Results

4.3.1 Gray-scale version

For this part, the results are not very good, we can see on the figure 4.6. We did 20 epoch for each experiment. We tried to do more epoch but the results getting worse and worse. We varied λ between 0.1 to 0.6 for the first network and we have the output image too warped when we use a small value for λ and not aligned. And when we use a bigger value the result is not much warped but not aligned too. With the second network, the output is too scale and not at all aligned with the RGB frame. The value of the loss function for each epoch is shown on the figure 4.7.




Figure 4.6: Results with gray-scale inputs for affine networks, voxelmorph network with $\lambda = 0.2$, $\lambda = 0.6$ respectively.




Figure 4.7: Loss function for affine network, voxelmorph network with $\lambda = 0.2$ and $\lambda = 0.6$ respectively.

4.3.2 Contours version

For this part of the experience, when we use the frames after to have applied some filters at the input of the networks, the resulting image does not move globally or really very little. With voxelmorph network, we have sometimes pixels that move locally (deform the shape) but never globally. We can see it on figure 4.8. We have the same kind of results with each contours filter. We tried to do more epoch but we get similar results.

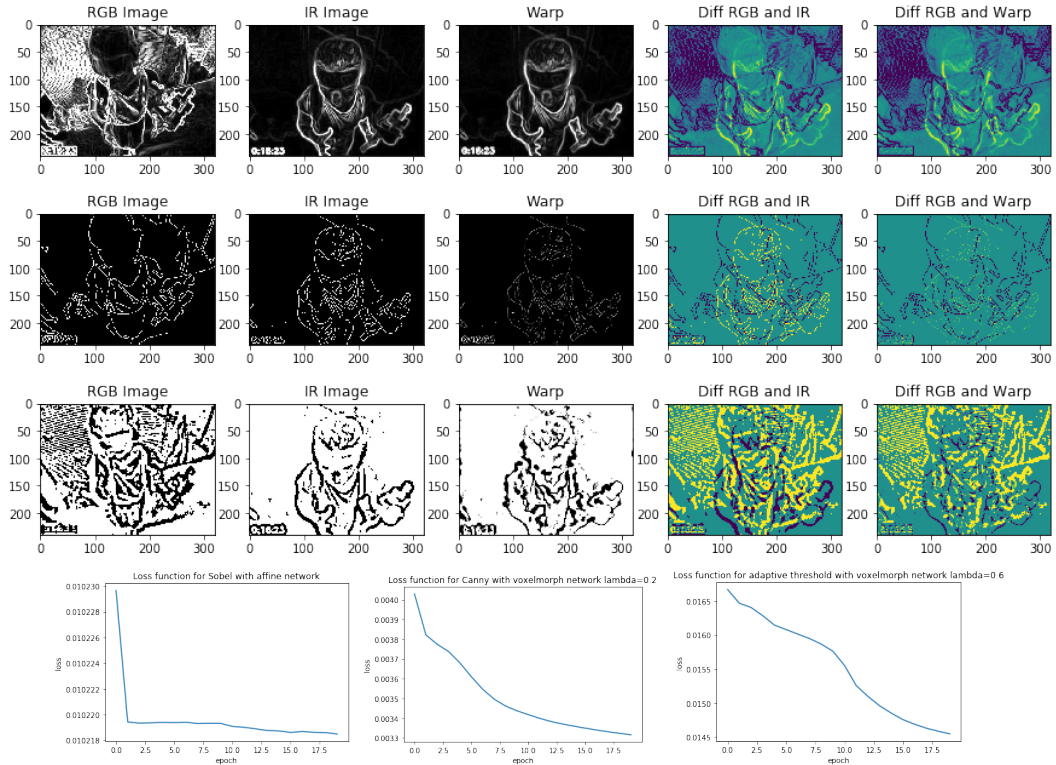


Figure 4.8: Results with contours inputs for affine networks with Sobel filter, voxelmorph network with $\lambda = 0.2$ with Canny filter and $\lambda = 0.6$ with adaptive threshold filter respectively. Below we have the loss function for each one.

4.3.3 Human segmentation version

For this part, we did 100 epoch and varied λ between 0.6 and 1. We have pretty good results and remark that the voxelmorph align better the frames but deform the shape in the image. So visually, the affine network give better results. Results are shown in figure 4.9. More results are available on the github page.

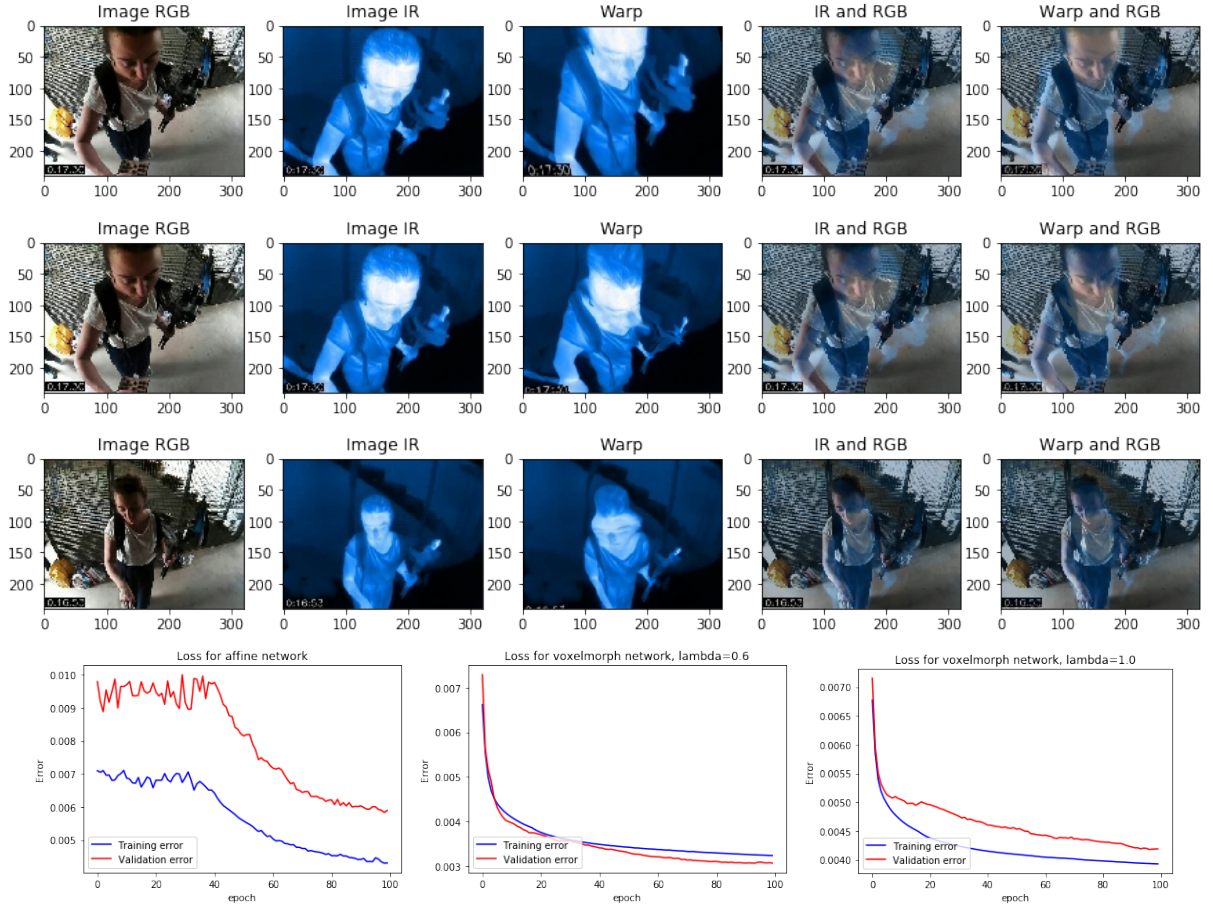


Figure 4.9: Results with humans segmentation inputs for affine networks, voxelmorph network with $\lambda = 0.6$, $\lambda = 1.0$ respectively. Below we have the loss function for each one.

Chapter 5

Discussion

In the first part of the experience the results are bad because we cannot really compare the RGB frame with the IR frame because the pixels value are not correlated between them. For example, if on the RGB image a person has a black t-shirt or a white t-shirt, on the IR image it will appear like it was the same ones. We can see that on figure 5.1. So if we do the mean square error between the IR and RGB images, it is going to measure the similitude between the two images that is not what we want. The only correlation between these two images, it is the shape borders in the image.

It is why in the second part of the experience we used the frames that contain the contours of the shapes. But we had not good results too. We thought it was because there is much more contrast on the RGB frames. We can see it on figure 5.2. In fact, we tried on very clean images, just images of contour of circles and we get the same results, the moving image does not move globally. So we understood that the problem was coming from the loss function. The problem is that the loss function is very flat (for the affine network). So if we move a little bit the image, the loss function will have nearly the same value or even the same value, thus the optimizer will not update the parameters of the network because the gradient will be very close to 0 or 0.

To counter this, instead of used the contours, we used the segmentation of the image. Because




Figure 5.1: We can see that the t-shirt white (very light) and the t-shirt black (very dark) on the RGB frame appear the same color on the IR frame




Figure 5.2: We remark there is so much contrast on the RGB image (on the left)




Figure 5.3: On the left, we can see an example of bad segmentation for the RGB image (on the left). On the right, the effect on the output of the first network when two frames are too far each others.

it was really complex to segment all the frames, we focused on the human segmentation of the frames. And as you can see on the figure 4.9, the results are really better, not perfect because the pretrain Deeplab network is not optimal to segment the humans as you can see on figure 5.3. To have better results, we must improve the segmentation of the images. The main problem is when two frames are too far each others, the case where the intersection between the two mask frames are empty. It is similar to the second part, when we move a little bit the frame, it will have the same value in the loss function, so the optimizer will not update the parameters or for the case of the voxelmorph network, it can do some weird results like the flatten of the image to minimize the difference between the two frames. We can see this kind of weird results in figure 5.3.

5.1 Future work

So for the next, if we want to have better results we must improve the segmentation step, whether for the RGB frames or the IR frames. In a first time, we must try to find a better pretrain deep network, if possible trained with infrared images too. But it is really complicated to find an already network trained.

The best way is to train a deep network with our data. But for that, we must have the ground truth label for a large number of frames. It asks a big work to mask a large numbers of frames




Figure 5.4: Example of a typical ground truth label of an image for semantic segmentation.

by hand. We can see an example of ground truth label on figure 5.4. With that, we can very well segment the images not only the humans but all the rest. And we will not have anymore problems when the two frames are too far each others because the entire image is segmented.

Chapter 6

Conclusion

As we could see, the registration between a IR and RGB image is not as obvious. We cannot just compare the similitude between the pixel values but we must find a correlation between the RGB and IR images. Without to use a correlation, an unsupervised deep-learning algorithm is not effective. The only correlation we can observe is the borders of the shapes in the image. So we applied some filters (Sobel, Canny , ...) to bring out these borders, but the result is not good yet because the loss function is too flat. So to overcome this problem, we used the human segmentation of the frames and we get much better results. Not again perfect, because the segmentation could be again better, whether the RGB frames segmentation (with Deeplab) or the IR frames segmentation (with thresh). The best way to optimize these segmentation is to train a deep network with our data but it requires to have the ground truth label for a large number of frames.

Bibliography

- [1] Unsupervised Learning of Probabilistic Diffeomorphic Registration for Images and Surfaces Adrian V. Dalca, Guha Balakrishnan, John Guttag, Mert R. Sabuncu MedIA: Medial Image Analysis. 2019. eprint arXiv:1903.03545
- [2] <https://en.wikipedia.org/wiki/Infrared>
- [3] Unsupervised Learning for Fast Probabilistic Diffeomorphic Registration Adrian V. Dalca, Guha Balakrishnan, John Guttag, Mert R. Sabuncu MICCAI 2018. eprint arXiv:1805.04605
- [4] VoxelMorph: A Learning Framework for Deformable Medical Image Registration Guha Balakrishnan, Amy Zhao, Mert R. Sabuncu, John Guttag, Adrian V. Dalca IEEE TMI: Transactions on Medical Imaging. 2019. eprint arXiv:1809.05231
- [5] An Unsupervised Learning Model for Deformable Medical Image Registration Guha Balakrishnan, Amy Zhao, Mert R. Sabuncu, John Guttag, Adrian V. Dalca CVPR 2018. eprint arXiv:1802.02604
- [6] U-Net: Convolutional Networks for Biomedical Image Segmentation Olaf Ronneberger, Philipp Fischer, Thomas Brox. eprint arXiv:1505.04597
- [7] Spatial Transformer Networks Max Jaderberg, Karen Simonyan, Andrew Zisserman, Koray Kavukcuoglu. eprint arXiv:1506.02025
- [8] Encoder-Decoder with Atrous Separable Convolution for Semantic Image Segmentation Liang-Chieh Chen and Yukun Zhu and George Papandreou and Florian Schroff and Hartwig Adam. ECCV. 2018
- [9] Searching for MobileNetV3 Andrew Howard and Mark Sandler and Grace Chu and Liang-Chieh Chen and Bo Chen and Mingxing Tan and Weijun Wang and Yukun Zhu and Ruoming Pang and Vijay Vasudevan and Quoc V. Le and Hartwig Adam. ICCV. 2019